

I've built a speech-to-speech conversion project that takes text from different file formats like documents or PDFs and turns it into voice. The system first extracts the text and gives me the option to generate audio using either an offline TTS engine or Google Text-to-Speech. Once the base audio is created, it is passed through Resemble AI's Speech-to-Speech mode, which clones the sound into my own voice. Right now, I'm testing the pipeline to check how natural the cloned voice sounds. The project is modular, with separate components for file handling, text extraction, TTS, and voice cloning, making it easy to extend or improve. In short, this assignment demonstrates how text and AI-powered speech cloning can be combined into a single working system.